

+ Code

+ Text

Start coding or generate with AI.

Linear Regression?

Linear Regression is a supervised machine learning algorithm used to predict a continuous value (number) based on input data.
It finds a straight line that best fits the data.

Linear Regression Equation

$y=mx+b$

Where:

y = Output (dependent variable)

x = Input (independent variable)

m = Slope (how fast y changes with x)

b = Intercept (value of y when x = 0)

Case Study: Predicting Student Marks

Problem Statement

A teacher wants to predict marks of a student based on the number of hours studied.



Dataset

Hours Studied (x)	Marks (y)
1	35
2	40
3	50
4	60
5	65

Observation

As study hours increase, marks also increase
Relationship is approximately linear

How Linear Regression Works (Step-by-Step)

- Plot data points on a graph
- Draw the best-fit straight line
- Minimize the error between actual and predicted values
- Use the line equation to predict future values

Example Prediction

Assume the trained model gives:
Predict marks for 4 hours of study
 $y=7(4)+30$
 $y=7(4)+30=58$

Predicted Marks = 58

Python Peactical

```
from sklearn.linear_model import LinearRegression
```

```
# Data
X = [[1], [2], [3], [4], [5]]
y = [35, 40, 50, 60, 65]

# Model
model = LinearRegression()
model.fit(X, y)

# Prediction
hours = [[4]]
predicted_marks = model.predict(hours)
print(predicted_marks)
```

[58.]

Advantages of Linear Regression

- Easy to understand and implement
- Works well for simple relationships
- Fast computation
- Useful for prediction and trend analysis

Limitations

- Assumes linear relationship
- Sensitive to outliers
- Not suitable for complex patterns

Real-Life Applications

- House price prediction
- Salary prediction
- Sales forecasting
- Crop yield estimation
- Weather trend analysis

Logistic Regression is a supervised machine learning algorithm used for binary classification problems.

It predicts the probability that an input belongs to a particular class (usually 0 or 1) using a sigmoid (S-shaped) function.

Simple Words

- Logistic Regression answers Yes/No, True/False, Pass/Fail
- Output is not a number like marks, but a probability between 0 and 1
- Based on this probability, the final class is decided

| Mathematical Form

$$P(y = 1) = \frac{1}{1 + e^{-(wx+b)}}$$

The sigmoid function converts any value into **0–1**

If probability $\geq 0.5 \rightarrow$ Class **1**

If probability $< 0.5 \rightarrow$ Class **0**

Start coding or [generate](#) with AI.

Example

Predict Pass (1) or Fail (0) based on study hours

Predict Disease (Yes/No) based on symptoms

Predict Spam (Yes/No) email

How Logistic Regression Works (Simple Explanation)

1. It takes the input feature (Hours Studied)
2. Applies a **linear equation**
3. Passes the result through a **Sigmoid function**
4. Produces a probability between 0 and 1

Sigmoid Function:

$$P(y = 1) = \frac{1}{1 + e^{-z}}$$

Where:

$$z = w \cdot x + b$$

Decision Rule

- If probability $\geq 0.5 \rightarrow$ Pass
- If probability $< 0.5 \rightarrow$ Fail

Example Prediction

Student studies for 3 hours

- Model output = 0.30
- Prediction $\rightarrow$  Fail

Student studies for 5 hours

- Model output = 0.85
- Prediction $\rightarrow$  Pass

Python Practical

```
from sklearn.linear_model import LogisticRegression

# Training data
X = [[1], [2], [3], [4], [5], [6]]
y = [0, 0, 0, 1, 1, 1]

# Model
model = LogisticRegression()
model.fit(X, y)

# Prediction
hours = [[5]]
prediction = model.predict(hours)
print("Pass" if prediction[0] == 1 else "Fail")
```

Pass

Advantages of Logistic Regression

Easy to understand
Fast to train
Works well for binary problems
Probability-based output

Limitations

Works best when classes are linearly separable
Not suitable for complex non-linear data

A hospital wants to predict whether a patient has a disease or not based on basic medical test results.

This is a binary classification problem:

Disease = 1
No Disease = 0

Dataset (Simple Example)

Age	Blood Sugar Level	Disease
25	120	0
30	130	0
40	150	1
50	170	1
60	180	1

Logistic Regression Works (Simple Steps)

Takes patient data (Age, Sugar Level)
Applies a linear equation
Uses sigmoid function
Outputs probability between 0 and 1
Classifies patient as Disease / No Disease

Q.1 predict whether a patient has a disease or not using multiple medical features with the help of Logistic Regression.

The model uses **four medical features**:

Feature	Description
Age	Patient age (years)
Blood Sugar	Sugar level (mg/dL)
Blood Pressure	BP (mmHg)
BMI	Body Mass Index

Logistic Regression?

Output has two classes only
Can handle multiple input features
Produces probability-based decision
Widely used in healthcare diagnosis

Difference Between Linear Regression and Logistic Regression

Linear Regression	Logistic Regression
Predicts a continuous value	Predicts a binary outcome
Output is a number	Output is 0 or 1 (Yes/No)

Output Type

Aspect	Linear	Logistic
Output Range	$-\infty$ to $+\infty$	0 to 1
Example	Marks = 78	Pass (1) / Fail (0)

Mathematical Model

Linear Regression

$y = mx + b$

Logistic Regression

$P(y = 1) = \frac{1}{1 + e^{-(mx+b)}}$

(Logistic regression uses a Sigmoid function)

Problem Type

Linear Regression	Logistic Regression
House price prediction	Email spam detection
Salary prediction	Disease (Yes/No)
Temperature forecasting	Pass / Fail result

Graph Shape

Linear Regression → Straight line

Logistic Regression → S-shaped (Sigmoid) curve

Learning Method

Feature	Linear	Logistic
Error Function	Mean Squared Error (MSE)	Log Loss (Cross-Entropy)
Decision Boundary	Not required	Required

Linear Regression predicts continuous values, whereas Logistic Regression predicts categorical (binary) outcomes using a sigmoid function.

Use Linear Regression → when output is a number

Use Logistic Regression → when output is Yes/No

Design a machine learning model using Logistic Regression to predict whether a crop is at risk of disease (Yes/No) based on environmental and crop-related features.

Feature	Description
Temperature (°C)	Average weekly temperature
Humidity (%)	Relative humidity
Rainfall (mm)	Weekly rainfall
Wind Speed (km/h)	Average wind speed
Crop Age (days)	Days after sowing

Output (Target Variable)

Value	Meaning
1	Disease Risk Present
0	No Disease Risk

Temp	Humidity	Rainfall	Wind	Crop Age	Disease
------	----------	----------	------	----------	---------

28	65	20	5	18	0
30	70	25	6	22	0
32	85	80	4	30	1
33	90	95	3	35	1
31	88	70	4	28	1

Collect weather and crop data

Apply a linear combination of features

Use sigmoid function

Obtain disease risk probability

Classify crop as Risk / No Risk

Early Prediction of Crop Disease Risk Using Machine Learning with Multi-Feature Environmental Data

Early Prediction of Water Quality and Contamination Risk Using Machine Learning

Background of the Problem

Access to clean and safe water is a major global challenge, particularly in developing countries.

Rivers, lakes, groundwater, and drinking water sources are increasingly contaminated due to:

Industrial discharge

Agricultural runoff (fertilizers and pesticides)

Urban sewage and wastewater

Poor monitoring infrastructure

Traditional water quality assessment relies on laboratory-based chemical analysis, which is:

Time-consuming

Costly

Not suitable for real-time monitoring

Difficult to scale across large regions

As a result, water contamination is often detected too late, leading to water-borne diseases, environmental damage, and public health crises.

Water Quality Checking Device Names (Real Devices)

Water quality can be monitored using devices such as a multiparameter water quality meter, water quality analyzer, portable testing kit, and IoT-based water quality monitoring systems.

Logistic Regression for Water-Borne Diseases?

Why Logistic Regression Fits Water-Borne Disease Problems

Water-borne disease prediction is usually a binary decision problem:

Disease Outbreak = Yes / No

Patient Infected = Yes / No

Water Safe = Safe / Unsafe

Logistic Regression Is Used (Conceptually)

Logistic Regression predicts the probability of disease occurrence based on multiple environmental, water quality, and demographic factors.

If probability $\geq$ threshold (e.g., 0.5) $\rightarrow$ High risk / Disease present

Input Features (Multiple, Real)

Water Quality Parameters

- pH
- Turbidity
- Total Dissolved Solids (TDS)
- Dissolved Oxygen (DO)
- Nitrate level
- Fecal coliform count

Environmental Factors

- Rainfall
- Temperature
- Flood events

Human & Demographic Factors

- Population density
- Sanitation level
- Water source type

Output Variable

Value	Meaning
1	Water-borne disease outbreak / infection
0	No outbreak / healthy population

Logistic Regression Model (Conceptual)

$$P(\text{Disease}) = \frac{1}{1 + e^{-(w_1pH + w_2Turb + w_3TDS + w_4Rain + w_5Pop + b)}}$$

Real Water-Borne Diseases Where Logistic Regression Is Used

1. Cholera

Common Symptoms (Features)

- Profuse watery diarrhea (rice-water stool)
- Severe dehydration
- Vomiting
- Muscle cramps
- Dry mouth and thirst
- Low blood pressure

2. Typhoid

Common Symptoms (Features)

- High fever (often >39°C)
- Headache
- Weakness and fatigue
- Abdominal pain
- Loss of appetite
- Constipation or diarrhea

3. Dysentery (डिसेंट्री)

Common Symptoms (Features)

- Bloody diarrhea
- Mucus in stool
- Abdominal cramps
- Fever

Nausea

Dehydration

3. Hepatitis A

Common Symptoms (Features)

Jaundice (yellow eyes/skin)

Dark urine

Pale stool

Fever

Fatigue

Loss of appetite

Nausea

4. Giardiasis(जियाडायसिस)

Common Symptoms (Features)

Chronic diarrhea

Greasy/foul-smelling stool

Abdominal bloating

Gas

Weight loss

Fatigue

Combined Feature Set (For ML / Logistic Regression)

Feature	Type
Fever	Binary / Continuous
Diarrhea	Binary / Frequency
Vomiting	Binary
Dehydration	Severity
Abdominal Pain	Severity
Jaundice	Binary
Stool Type	Categorical
Fatigue	Binary
Appetite Loss	Binary

Target Variable (Output)

Value Meaning

1 Water-borne disease present

0 No water-borne disease

In Logistic Regression for water-borne diseases:

Symptoms = independent variables (features)

Disease presence = dependent variable (Yes/No)

Step 1: Install Required Libraries

`pip install numpy pandas scikit-learn shap matplotlib`

Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (2.0.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (2.2.2)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.12/dist-packages (1.6.1)
Requirement already satisfied: shap in /usr/local/lib/python3.12/dist-packages (0.50.0)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (3.10.0)
Requirement already satisfied: python-dateutil<=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.3)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.16.3)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.5.3)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.12/dist-packages (from shap) (4.67.1)
Requirement already satisfied: packaging>20.0 in /usr/local/lib/python3.12/dist-packages (from shap) (25.0)
Requirement already satisfied: slicer==0.0.8 in /usr/local/lib/python3.12/dist-packages (from shap) (0.0.8)
Requirement already satisfied: numba>=0.54 in /usr/local/lib/python3.12/dist-packages (from shap) (0.60.0)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.12/dist-packages (from shap) (3.1.2)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.12/dist-packages (from shap) (4.15.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (4.61.1)

```
Requirement already satisfied: Kivisoiver>=1.5.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.4.9)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (3.3.1)
Requirement already satisfied: llvmlite<0.44, >=0.43.0dev0 in /usr/local/lib/python3.12/dist-packages (from numba>=0.54->shap) (0.43.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
```

Step 2: Import Libraries

```
import numpy as np
import pandas as pd
import shap
import matplotlib.pyplot as plt

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

Step 3: Create Sample Water Quality Dataset

```
data = {
    "pH": [6.8, 7.1, 6.2, 5.9, 7.4, 6.0, 6.5, 5.8],
    "Turbidity": [3, 2, 8, 10, 1, 9, 6, 11],
    "TDS": [300, 280, 700, 820, 250, 760, 680, 900],
    "DO": [6.5, 7.0, 3.8, 3.2, 7.2, 3.5, 4.0, 2.8],
    "Nitrate": [20, 18, 45, 60, 15, 55, 48, 70],
    "Disease_Risk": [0, 0, 1, 1, 0, 1, 1, 1]
}

df = pd.DataFrame(data)
```

Step 4: Split Features and Target

```
X = df.drop("Disease_Risk", axis=1)
y = df["Disease_Risk"]
```

Step 5: Train-Test Split

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, random_state=42
)
```

Step 6: Feature Scaling

```
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Step 7: Train Logistic Regression Model

```
model = LogisticRegression()
model.fit(X_train_scaled, y_train)
```

```
LogisticRegression
```

Step 8: Make Prediction (Example)

```
sample = pd.DataFrame({
    "pH": [6.1],
    "Turbidity": [9],
    "TDS": [780],
    "DO": [3.2],
    "Nitrate": [58]
})

sample_scaled = scaler.transform(sample)
prediction = model.predict(sample_scaled)
probability = model.predict_proba(sample_scaled)

print("Prediction:", "High Risk" if prediction[0] == 1 else "Low Risk")
print("Probability:", probability)
```


```
Prediction: High Risk
Probability: [[0.04670799 0.95329201]]
```

Step 9: Apply XAI Using SHAP


```
explainer = shap.Explainer(model, X_train_scaled)
shap_values = explainer(X_test_scaled)
```

Step 10: SHAP Feature Importance (Global Explanation)

```
# shap.plots.waterfall(shap_values[0], feature_names=X.columns)
shap.plots.waterfall(shap_values[0])
```



```
shap.summary_plot(shap_values, X_test, feature_names=X.columns)
```




SHAP Bar Plot (Mean Importance)

Average contribution of each feature

Clean and simple (journal-friendly)

```
shap.plots.bar(shap_values)
```



3 SHAP Dependence Plot (Feature Interaction)


What it shows

How one feature affects prediction

Relationship between feature value and risk

```
shap.dependence_plot(
    "Turbidity", shap_values.values, X_test
)
```





4. SHAP Force Plot (Single Prediction – Interactive)

What it shows

Push & pull of features for one sample

Best for presentations and demos

```
shap.initjs()
shap.force_plot(
    shap_values.base_values[0],
    shap_values.values[0],
    X_test.iloc[0]
)
```



Start coding or [generate](#) with AI.

What is SVM?

Support Vector Machine (SVM) is a supervised machine learning algorithm used for classification and regression.

It works by finding an optimal decision boundary (hyperplane) that maximizes the margin between different classes.

Intuition (Simple Words)

SVM draws a line (2D) or plane (3D) or hyperplane (high-D)

The best boundary is the one with the maximum distance from the nearest points of both classes

These nearest points are called support vectors

Mathematical Idea (Classification)

For binary classification, SVM solves:

$$\min_{w,b} \frac{1}{2} \|w\|^2 \quad \text{subject to } y_i(w \cdot x_i + b) \geq 1$$

- w : weight vector (orientation of hyperplane)
- b : bias
- x_i : data point
- $y_i \in \{-1, +1\}$: class label

👉 This **maximizes the margin**.

What are Support Vectors?

Data points closest to the decision boundary

They define the position of the hyperplane

Removing other points does not change the boundary

Kernels (Key Strength of SVM)

When data is not linearly separable, SVM uses kernels to map data to a higher dimension.

Kernel	Use Case
Linear	Linearly separable data

Polynomial	Curved boundaries
RBF (Gaussian)	Complex patterns (most popular)
Sigmoid	Neural-network-like behavior

Simple Example

Problem: Classify emails as Spam / Not Spam

Features: word frequency, links, sender score

SVM: finds a boundary that best separates spam vs non-spam with maximum margin.

```
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_breast_cancer

X, y = load_breast_cancer(return_X_y=True)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

model = SVC(kernel="rbf", C=1.0, gamma="scale")
model.fit(X_train, y_train)

accuracy = model.score(X_test, y_test)
print("Accuracy:", accuracy)
```

Accuracy: 0.9473684210526315


How to Change Kernel in Support Vector Machine (SVM)

In scikit-learn, you change the kernel by modifying the kernel parameter in SVC().

1 Linear Kernel (for linearly separable data)

```
from sklearn.svm import SVC

model = SVC(kernel="linear", C=1.0)
model.fit(X_train, y_train)
```



```
SVC(kernel='linear')
```

```
y_pred = model.predict(X_test)
```

```
from sklearn.metrics import accuracy_score

accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.956148350877193

One-Line Shortcut (Also Correct)

```
accuracy = model.score(X_test, y_test)
print("Accuracy:", accuracy)
```

Accuracy: 0.956148350877193

Accuracy

$$\text{Accuracy} = \frac{\text{Correct Predictions}}{\text{Total Predictions}}$$

Example:

- Correct predictions = 90
- Total samples = 100

$$\text{Accuracy} = 90/100 = 0.90 = 90\%$$

```
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
```



```
# Train model
model = SVC(kernel="linear", C=1.0)
model.fit(X_train, y_train)

# Predict
y_pred = model.predict(X_test)

# Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
```

Accuracy: 0.96

Naive Bayes

Naive Bayes is a supervised machine learning algorithm based on Bayes' Theorem.

It is mainly used for classification problems, especially text and categorical data.

It is called "Naive" because it assumes that all features are independent of each other.

Simple Intuition

Uses probability to predict a class

Calculates the probability of each class

Chooses the class with the highest probability

Bayes' Theorem

 Bayes' Theorem

$$P(C|X) = \frac{P(X|C) P(C)}{P(X)}$$

Where:

- $P(C|X)$: Posterior probability (class after seeing data)
- $P(X|C)$: Likelihood
- $P(C)$: Prior probability
- $P(X)$: Evidence

Types of Naive Bayes

Type	Used When
Gaussian NB	Continuous data (marks, height)
Multinomial NB	Text data (word counts)
Bernoulli NB	Binary features (Yes/No)

Problem

Predict whether a student **Pass / Fail** based on **Study Hours**.

Hours Studied	Result
Low	Fail
Medium	Pass
High	Pass

Naive Bayes computes:

- Probability of **Pass**
- Probability of **Fail**
- Chooses the class with **higher probability**

```
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import pandas as pd

# Sample data
data = {
    "Hours": [1, 2, 3, 4, 5, 6],
    "Result": [0, 0, 0, 1, 1, 1]
}
```

```

df = pd.DataFrame(data)

X = df[["Hours"]]
y = df["Result"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)

model = GaussianNB()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))

```

Accuracy: 0.0

Advantages of Naive Bayes

Very fast

Works well with text data

Requires less training data

Easy to implement

Where Naive Bayes is Used

Spam email detection

Sentiment analysis

Document classification

Medical diagnosis (baseline model)

Aspect	Naive Bayes	SVM
Speed	Very fast	Slower
Data size	Small-medium	Medium-large
Interpretability	Medium	Low
Text data	Excellent	Moderate

Problem Statement

A college wants to predict whether a student will Pass or Fail based on simple academic features using the Naive Bayes algorithm.

This is a binary classification problem:

Pass = 1

Fail = 0

Dataset (Simple & Realistic)

Study Hours	Attendance (%)	Internal Marks	Result
Low	60	40	Fail
Medium	70	55	Fail
Medium	80	65	Pass
High	85	75	Pass
High	90	80	Pass

Features Used (Input Data)

Feature	Description
Study Hours	Hours studied per day
Attendance	Percentage of attendance

Why Naive Bayes?

Works well with small datasets
 Easy to understand
 Fast computation
 Suitable for categorical data
 Good baseline classifier

How Naive Bayes Works (Simple Steps)

Calculate probability of Pass and Fail
 Calculate probability of each feature for each class
 Apply Bayes' Theorem
 Choose the class with highest probability

Bayes' Theorem Used

$$P(Class|Features) = \frac{P(Features|Class) \times P(Class)}{P(Features)}$$

Naive assumption:

All features are **independent** of each other

Example Prediction

New Student Details:

- Study Hours = Medium
- Attendance = 75%
- Internal Marks = 60

Naive Bayes Output:

- $P(\text{Pass}) = 0.78$
- $P(\text{Fail}) = 0.22$

✓ Prediction: PASS

Double-click (or enter) to edit

```
import pandas as pd
from sklearn.naive_bayes import GaussianNB
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Dataset
data = {
    "StudyHours": ["Low", "Medium", "Medium", "High", "High"],
    "Attendance": [60, 70, 80, 85, 90],
    "InternalMarks": [40, 55, 65, 75, 80],
    "Result": ["Fail", "Fail", "Pass", "Pass", "Pass"]
}

df = pd.DataFrame(data)

# Encode categorical data
le = LabelEncoder()
df["StudyHours"] = le.fit_transform(df["StudyHours"])
df["Result"] = le.fit_transform(df["Result"])

X = df[["StudyHours", "Attendance", "InternalMarks"]]
y = df["Result"]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)

# Naive Bayes model
```

```
model = GaussianNB()
model.fit(X_train, y_train)

# Prediction
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Accuracy: 0.5

Metrics Implementation for Naive Bayes (NB)

```
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    fbeta_score,
    classification_report
)
```

Step 2: Predict Using the Trained NB Model

(Using the model you already trained)

Double-click (or enter) to edit

```
y_pred = model.predict(X_test)
```

Step 3: Calculate Evaluation Metrics

Accuracy

```
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.5

Precision

```
precision = precision_score(y_test, y_pred)
print("Precision:", precision)
```

Precision: 0.5

Recall

```
recall = recall_score(y_test, y_pred)
print("Recall:", recall)
```

Recall: 1.0

F1-score

```
f1 = f1_score(y_test, y_pred)
print("F1-score:", f1)
```

F1-score: 0.6666666666666666

All Metrics Together (Recommended)

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.50	1.00	0.67	1
accuracy			0.50	2
macro avg	0.25	0.50	0.33	2
weighted avg	0.25	0.50	0.33	2

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control the handling of labels with no predicted samples. See https://scikit-learn.org/stable/user_guide.html#handling-of-undefined-values for details.

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control the handling of labels with no predicted samples. See https://scikit-learn.org/stable/user_guide.html#handling-of-undefined-values for details.

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control the handling of labels with no predicted samples. See https://scikit-learn.org/stable/user_guide.html#handling-of-undefined-values for details.

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control the handling of labels with no predicted samples. See https://scikit-learn.org/stable/user_guide.html#handling-of-undefined-values for details.

Metric

Meaning

Accuracy

Overall correctness of the model

Metric	What it measures on the model
Precision	How many predicted Pass are actually Pass
Recall	How many actual Pass students are detected
F1-score	Balance between Precision & Recall
F2-score	Recall-weighted score (missed Pass is costly)

Accuracy → overall correctness

Precision → correctness of positive predictions

Recall → ability to find all positives

F1 → balance of precision & recall

F2 → recall-focused metric

Decision Tree

Decision Tree is a supervised machine learning algorithm used for classification and regression.

It works like a flowchart, where:

Internal nodes → questions (conditions)

Branches → answers (Yes/No)

Leaf nodes → final decision/output

Simple Intuition

Think of how we make decisions in real life:

If marks ≥ 40 → Pass

Else → Fail

A Decision Tree learns such if-else rules automatically from data.

Types of Decision Trees

Classification Tree → output is a class (Pass/Fail, Yes/No)

Regression Tree → output is a number (marks, price)

How Decision Tree Works

Splits data based on the best feature

Uses measures like:

Gini Index

Entropy

Information Gain

Continues splitting until a stopping condition is met

Problem

Predict whether a student will Pass or Fail based on:

Study Hours

Attendance

Study Hours	Attendance (%)	Result
Low	60	Fail
Medium	70	Fail
Medium	80	Pass
High	85	Pass
High	90	Pass

Decision Rule (Learned by Tree)

If Attendance ≥ 75:

Pass

Else:

Fail

Simple Python Code

```
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score

# Dataset
data = {
    "StudyHours": ["Low", "Medium", "Medium", "High", "High"],
    "Attendance": [60, 70, 80, 85, 90],
    "Result": ["Fail", "Fail", "Pass", "Pass", "Pass"]
}

df = pd.DataFrame(data)

# Encode categorical feature
le = LabelEncoder()
df["StudyHours"] = le.fit_transform(df["StudyHours"])
df["Result"] = le.fit_transform(df["Result"])

X = df[["StudyHours", "Attendance"]]
y = df["Result"]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)

# Decision Tree model
model = DecisionTreeClassifier(criterion="gini")
model.fit(X_train, y_train)

# Prediction
y_pred = model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
```

Accuracy: 1.0

Advantages of Decision Tree

Easy to understand and interpret

No need for feature scaling

Works with categorical and numerical data

Visualizable

Limitations

Can overfit the data

Small change in data → different tree

Less stable than ensemble method

Algorithm	Interpretability	Overfitting
Naive Bayes	Medium	Low
Decision Tree	High	High
Random Forest	Medium	Low

Supervised Learning (Basic Idea)

Supervised Learning is a learning method where:

Input data (X) and correct output (Y) are known

The model learns a mapping $X \rightarrow Y$

Error is calculated and minimized

Examples: Pass/Fail, Disease Yes/No, Spam/Not Spam

perceptron Learning

Definition

The Perceptron is the simplest supervised learning algorithm and the basic unit of a neural network.

Key Points

Binary classifier (0/1)

Uses step (threshold) activation

Works only for linearly separable data

Limitation

Cannot solve XOR problem

Single-Layer Neural Network

Definition

Has only one layer of weights

No hidden layer

Example: Perceptron, ADALINE

Capability

Solves simple linear problems

Cannot solve non-linear problems

✖ Multilayer Neural Network (MLP)

Definition

Has input layer, hidden layer(s), output layer

Uses non-linear activation functions

Advantage

Can solve complex and non-linear problems

Solves XOR problem

Double-click (or enter) to edit

✖ ADALINE (Adaptive Linear Neuron)

Definition

ADALINE is an improved Perceptron that:

Uses linear activation

Minimizes Mean Squared Error (MSE)

Uses gradient descent

Difference from Perceptron

Double-click (or enter) to edit

✖ Linear Separability

Meaning

Data is linearly separable if classes can be separated by a straight line (or plane).

◆ Examples

Problem	Linearly Separable
AND, OR	✅ Yes
XOR	❌ No

✖ Perceptron & ADALINE fail on XOR

Double-click (or enter) to edit

✖ ADALINE (Adaptive Linear Neuron)

Definition

ADALINE is an improved Perceptron that:

- Uses linear activation
- Minimizes Mean Squared Error (MSE)
- Uses gradient descent
- Difference from Perceptron

Perceptron	ADALINE
Step function	Linear function
Discrete error	Continuous error
Less stable	More stable

✖ Still fails on XOR

MADALINE (Multiple ADALINE)

- Definition
- MADALINE is a network of multiple ADALINE units.
- First multilayer neural network
- Uses majority voting
- Capability
- More powerful than ADALINE
- Handles moderately complex patterns

Backpropagation Network (BPN)

- Definition
- Backpropagation is a learning algorithm used to train multilayer neural networks.
- Working Steps
- Forward pass
- Error calculation
- Backward pass
- Weight update
- Advantage
- Solves non-linear problems
- Foundation of deep learning

⌵ RBFN (Radial Basis Function Network)

- Definition
- RBFN is a neural network that uses radial basis functions in the hidden layer.
- Structure
- Input layer
- RBF hidden layer
- Output layer
- Key Feature
- Neurons respond based on distance from center
- Applications
- Pattern recognition
- Function approximation

Model	Linear / Non-Linear	Hidden Layer
Perceptron	Linear	No
ADALINE	Linear	No
MADALINE	Semi-nonlinear	Yes
MLP + Backprop	Non-linear	Yes
RBFN	Non-linear	Yes

Perceptron: A basic supervised learning algorithm for binary classification.

ADALINE: An adaptive linear neuron using MSE and gradient descent.

MADALINE: A multilayer network composed of multiple ADALINE units.

Backpropagation: An algorithm that updates weights by propagating error backward.

RBFN: A neural network using radial basis functions in hidden layers.

Perceptron is the simplest supervised learning algorithm and the basic building block of neural networks.

It is used for binary classification, where the output is 0 or 1 (Yes/No, Pass/Fail).

In Simple Words

Perceptron works like a decision maker

It takes inputs, multiplies them by weights

Adds them together

Applies an activation function

Gives a binary output

Mathematical Form

$$y = \begin{cases} 1, & \text{if } w \cdot x + b > 0 \\ 0, & \text{otherwise} \end{cases}$$

Where:

- x = input values
- w = weights
- b = bias

Key Points for Students

Works only for linearly separable data

Cannot solve XOR problem

Used for binary classification

Foundation of modern neural networks

Problem

Predict whether a student Pass (1) or Fail (0) based on study hours.

▼ Create a Small Dataset

```
[1] ✓ Os
import numpy as np

# Input feature: study hours
X = np.array([[1], [2], [3], [4], [5]])

# Target: 0 = Fail, 1 = Pass
y = np.array([0, 0, 0, 1, 1])
```

▼ Step 2: Train a Perceptron Model

```
[2] ✓ 4s
from sklearn.linear_model import Perceptron

# Create Perceptron model
model = Perceptron(max_iter=1000, eta0=0.1, random_state=42)

# Train the model
model.fit(X, y)
```

Perceptron
Perceptron(eta0=0.1, random_state=42)

▼ Step 3: Make Predictions

```
[3] ✓ Os
# Predict for a new student (studies 4 hours)
new_student = np.array([[4]])
```

```
prediction = model.predict(new_student)

print("Prediction:", "Pass" if prediction[0] == 1 else "Fail")
```

Prediction: Pass

Step 4: Check Accuracy (Optional but Good Practice)

```
from sklearn.metrics import accuracy_score

y_pred = model.predict(X)
accuracy = accuracy_score(y, y_pred)

print("Accuracy:", accuracy)
```

Accuracy: 1.0

Perceptron on XOR Problem (Failure Demonstration)

Step 1: XOR Dataset XOR Truth Table

X1	X2	Output
0	0	0
0	1	1
1	0	1
1	1	0

Step 2: Create XOR Dataset in Python

```
import numpy as np

# XOR inputs
X = np.array([
    [0, 0],
    [0, 1],
    [1, 0],
    [1, 1]
])

# XOR outputs
y = np.array([0, 1, 1, 0])
```

Step 3: Train a Perceptron Model

```
from sklearn.linear_model import Perceptron

model = Perceptron(max_iter=1000, eta0=0.1, random_state=42)
model.fit(X, y)
```

Perceptron
Perceptron(eta0=0.1, random_state=42)

Step 4: Predict on XOR Data

```
y_pred = model.predict(X)

print("Predictions:", y_pred)
print("Actual      :", y)
```

Predictions: [0 0 0 0]
Actual : [0 1 1 0]

Step 5: Check Accuracy

```
from sklearn.metrics import accuracy_score

accuracy = accuracy_score(y, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.5

Reason 1: Linear Model

Perceptron decision function:

Reason 2: XOR is NOT Linearly Separable

Points (0,1) and (1,0) → Class 1

Points (0,0) and (1,1) → Class 0

No single straight line can separate them

Visual Explanation (Exam Tip)

AND, OR → ✓ Linearly separable

XOR → ✗ Not linearly separable

- ADALINE (Adaptive Linear Neuron) is a supervised learning algorithm and an improved version of the Perceptron.

It uses a linear activation function and learns by minimizing the Mean Squared Error (MSE) using gradient descent.

ADALINE is a single-layer neural network

It compares the actual output with linear output

Updates weights smoothly using error

More stable and accurate than Perceptron

- Mathematical Form

$$y = w \cdot x + b$$

Error:

$$\text{Error} = (y_{\text{actual}} - y_{\text{predicted}})$$

Cost Function (MSE):

$$E = \frac{1}{2} \sum (y - \hat{y})^2$$

ADALINE (Practical)

ADALINE is NOT directly available in sklearn, so we implement it manually.

- Step 2B: ADALINE Implementation

```
[141] ✓ Os
import numpy as np

class Adaline:
    def __init__(self, lr=0.01, epochs=1000):
        self.lr = lr
        self.epochs = epochs

    def fit(self, X, y):
        # X shape: (n_samples, n_features)
        n_samples, n_features = X.shape

        self.w = np.zeros(n_features)
        self.b = 0.0

        for _ in range(self.epochs):
            # Linear output (n_samples,)
            linear_output = np.dot(X, self.w) + self.b

            # Error (n_samples,)
            errors = y - linear_output

            # Gradient descent update
            self.w += self.lr * np.dot(X.T, errors)
            self.b += self.lr * errors.sum()

    def predict(self, X):
```

```
linear_output = np.dot(X, self.w) + self.b
return np.where(linear_output >= 0.5, 1, 0)
```

Step 3B: Train & Predict Using ADALINE

```
[13]: In [13]: adaline = Adaline(lr=0.01, epochs=1000)
        adaline.fit(X, y)

        a_pred = adaline.predict(X)

        print("ADALINE Predictions:", a_pred)
```

ADALINE Predictions: [0 0 0 1]

This means ADALINE predicted the following classes for your input samples:

Sample	Actual Output (y)	ADALINE Prediction
1	0	0 ✓
2	0	0 ✓
3	0	0 ✓
4	1	1 ✓

Why Did ADALINE Give This Output?

Reason 1: Data is Linearly Separable

our dataset (study hours → pass/fail) can be separated by a straight line, so ADALINE can learn it.

Reason 2: Linear Activation + Gradient Descent

ADALINE:

Computes a linear output

Minimizes Mean Squared Error (MSE)

Learns smoothly, not abruptly like Perceptron

ADALINE predicted the correct class labels because the data is linearly separable and the model learned optimal weights.

Colab paid products - Cancel contracts here

Variables Terminal



✓ 9:24 AM Python 3